

The Rise of Verbal Reinforcement Learning: A Survey

Anonymous ACL submission

Abstract

Natural language is emerging as a primary feedback channel for improving language agents, capable of conveying intent, preferences, and causal structure in forms interpretable by both humans and modern language models. We call this paradigm Verbal Reinforcement Learning (VRL) and offer the first unified account of it, **with explicit inclusion criteria that separate it from RLHF, instruction following, and language-conditioned RL**. We organize the field around a single axis, *when* verbal feedback takes effect in an agent’s lifecycle and *what* it modifies, yielding three pillars: (1) **Language as Grounding Signal**, where language defines the task itself by specifying goals, states, and reward structures; (2) **Language as Deliberative Feedback**, where natural language guides reasoning at test time without the need to update model parameters; (3) **Language as Learning Signal**, where language-based feedback shapes model parameters through training. Within each pillar, we synthesize representative work, distinguish key subcategories of approaches, and outline the distinct role language plays in shaping agent behavior. Together, this taxonomy shows how verbal reinforcement is reshaping agent development, while also defining the challenges and opportunities for building more capable and aligned agents.

1 Introduction

Classical reinforcement learning has achieved remarkable successes by coupling agents with well-specified environments, action spaces, reasoning mechanisms, and task-specific reward functions. This framework has enabled agents to master Atari games (Mnih et al., 2015), defeat world champions at Go (Silver et al., 2016), optimize chip placement (Mirhoseini et al., 2021), and control complex physical systems (Degraeve et al., 2022). However, each success rests on heavy domain-specific engineering: the environment must be formalized, the




Figure 1: Language plays three complementary roles in agent’s development lifecycle: grounding the environment, scaffolding intermediate reasoning, and providing learning signal for further training.

interaction space defined, and the reward function must closely capture the intended behavior (Vamplew et al., 2021). Designing such specifications is a long-standing bottleneck (Amodi et al., 2016), especially in settings where tasks and goals are ambiguous, context-dependent, and difficult to encode directly.

The emergence of large language models (Brown et al., 2020) has introduced an alternative supervisory channel: natural language. Instead of requiring all task intent to be predefined, verbal feedback enables on-the-fly communication of corrections, preferences, and constraints in a manner that agents are increasingly able to interpret and implement. Modern large language models, particularly frontier models (Achiam et al., 2023), can accurately describe a wide range of phenomena. These capabilities include describing environments (Wang et al., 2023a), identifying bugs in code (McAleese et al., 2024), critiquing generated outputs (Madaan et al., 2023), and translating human preferences into actionable training signals (Hong et al., 2024a; Meng et al., 2024). Similar to how self-supervised learning (Devlin et al., 2019; Chen et al., 2020) expanded the use of unlabeled data in model development, verbal feedback is emerging as an important source of supervision for language-model-based




Figure 2: Our proposed taxonomy of Verbal Reinforcement Learning. Language serves three distinct functional roles: *grounding signal*, which structures environment interaction; *deliberative feedback*, which operates at inference time to refine reasoning; and *learning signal*, which drives parameter updates at training time.

agents.

We refer to this family of methods as **Verbal Reinforcement Learning (VRL)**, and we define it by two inclusion criteria: (i) *the feedback is expressed in natural language, whether self-generated, provided by a human, or produced by an external tool or model, and its linguistic structure is functionally consumed by the agent, rather than being collapsed to a scalar or preference bit before use; and (ii) the feedback participates in an improvement loop, that is, an output revision, a memory write, or a parameter update*. Both criteria are required, and they fence VRL off from adjacent paradigms. Plain instruction following satisfies neither, since no improvement loop exists. Vanilla RLHF and DPO satisfy only the second, because the verbal judgment is reduced to a preference bit before learning; we discuss them only as the scalar endpoint of the compression spectrum in Section 5. Language-conditioned RL treats language as an input to the policy rather than as feedback on its behavior. One boundary case the criteria settle: language consumed at specification time to construct a reward function, as in reward-code generation, satisfies criterion (i) even though the constructed function later emits scalars. The criteria also fix the modality scope: the feedback signal is natural-language text, while the agent’s observations may be non-textual, as in embodied settings where verbal corrections accompany visual state. The term was introduced by Shinn et al. (2023) to describe self-reflective agents (specific instance of what we call deliberative feedback); we deliberately generalize it to any method meeting the two criteria above, and we state this generalization explicitly to avoid conflation with the original, narrower sense.

Early results demonstrate that verbal feedback can drive substantial improvements across diverse settings. Ahn et al. (2022) used language to ground high-level instructions in physical affordances, enabling real-world robotic planning. Shinn et al. (2023) achieved 91% pass@1 on HumanEval through verbal self-reflection alone, and Madaan et al. (2023) showed that iterative verbal critique yields substantial gains across diverse generation tasks. Ouyang et al. (2022) demonstrated that a 1.3B-parameter model trained on verbal preference judgments outperformed the 175B GPT-3 baseline, a $130\times$ size disadvantage overcome by richer supervision. These are not isolated demonstrations and publications on arXiv referencing “self-refine,” “self-correction,” “verbal feedback,” and “language-feedback alignment” have grown from a handful in early 2020 to several hundred today, spanning multiple domains including coding assistants (Yang et al., 2024; Jimenez et al., 2024), scientific discovery (Bran et al., 2024; Lu et al., 2024), robotics (Wang et al., 2023a), mathematical reasoning (Hatamizadeh et al., 2026; Liu et al., 2026b; Shi et al., 2025a), clinical decision-making (Huang et al., 2025b; Zhou et al., 2025b) and education (Qian et al., 2025; Ahtisham et al., 2026).

These methods differ in their mechanisms (Figure 1): some operate during inference, others modify training data, and some redefine the task itself. However, they share a common characteristic: the use of natural language to guide, refine, and ground agent behavior. Yet no existing framework unifies them. Prior surveys have examined language-conditioned policies in text-based environments (Luketina et al., 2019), the bidirectional synergies between LLMs and RL (Pternea et al.,

2024), LLM self-correction in isolation (Kamoi et al., 2024), automatic correction strategies (Pan et al., 2024), feedback integration for text generation (Fernandes et al., 2023), preference optimization (Kaufmann et al., 2023; Xiao et al., 2024), agent memory (Zhang et al., 2025e), test-time scaling (Zhang et al., 2025b), self-evolving agents (Gao et al., 2025; Fang et al., 2025), and reinforcement learning for LLM agents (Zhang et al., 2025a); Natural Language Reinforcement Learning (Feng et al., 2024) recasts RL components in language space but is a framework proposal rather than a survey. Table 3 in Appendix A positions our taxonomy against these works: they organize by mechanism or by agent capability, whereas we organize by the signal itself. In contrast, our survey focuses on the functional role of *verbal feedback* as a first-class signal for improving agent, organized along a single axis: *when* natural language enters the agent’s lifecycle and *what* it modifies. We formalize this axis into a three-pillar taxonomy in Section 2. This paper makes following contributions:

- We present the first comprehensive survey of Verbal Reinforcement Learning, **defining the paradigm through explicit inclusion criteria that separate it from RLHF, instruction following, and language-conditioned RL**, and proposing a three-pillar taxonomy (Section 2) that organizes the field by *when* language acts in the agent lifecycle.
- Within each pillar, we synthesize representative methods, identify cross-cutting challenges and outline future directions for developing robust, capable and aligned Verbal RL agents.
- **We give explicit decision rules for methods whose signals span pillars, including a hybrid-cases table (Table 1), so the taxonomy is applicable rather than merely descriptive.**
- **We distill practitioner guidance for choosing among verbal-feedback mechanisms (Appendix B) and position the survey against prior overviews (Appendix A).**

We organize the paper as follows. Section 2 introduces the three-pillar taxonomy and a working example showing how the pillars operate together in a single coding agent. Sections 3–5 develop each pillar of the taxonomy in detail: language

as grounding signal, as deliberative feedback, and as learning signal, respectively. Finally, section 6 presents cross-cutting insights and outlines future research directions, while section 7 provides the conclusion.

2 Taxonomy

We organize VRL methods along a single axis: *when* natural language enters the agent’s lifecycle and, consequently, *what* it modifies. This yields three pillars, summarized in Figure 2 and synthesized in Table 2.

- **Pillar 1: Language as Grounding Signal (§3)** acts at *problem-definition time*, specifying the MDP (Sutton and Barto, 1998) itself: goals, states, actions, and rewards. A natural-language task description such as “*You are in a room with two doors*” defines what the agent perceives, how it can act, and what counts as success, all before any policy runs.
- **Pillar 2: Language as Deliberative Feedback (§4)** acts at *inference time*, refining outputs or reasoning traces through critique, memory, debate, or search, without updating parameters. A message such as “*Step 3 has an error; try dividing by 2, not 4*” redirects a single generation episode but leaves the model unchanged for future tasks.
- **Pillar 3: Language as Learning Signal (§5)** acts at *training time*, where verbal feedback is distilled into gradient updates that persistently reshape the policy. A judgment such as “*Response A is better than B because it addresses the edge case*” becomes a preference pair that alters model weights, affecting all subsequent interactions.

The key distinction is *persistence*: grounding defines the task space, deliberative feedback improves a single episode, and learning signal permanently reshapes the policy. We discuss each pillar in consecutive sections. **For methods whose signals span pillars, we use an explicit assignment rule: a method is placed by what its linguistic signal modifies at the moment the signal takes effect.** Table 1 applies this rule to the recurring hybrid cases and records their secondary roles instead of forcing each method into a single cell. Throughout, MDP vocabulary serves as an organizing lens rather than a claim that every method admits an MDP formulation; many language-agent settings are partially

observable, non-Markovian, or lack explicit transition and reward functions, and we note where the formalism stops applying.

Table 1: Assignment rule applied to hybrid cases: placement follows what the signal modifies at the moment it takes effect; secondary roles are tagged, not suppressed.

Hybrid case	Primary	Secondary role
Reward-code generation (Ma et al., 2024)	P1	Constructed reward later trains a policy (P3)
Experiential memory (Shinn et al., 2023)	P2	Lessons persist across episodes, bordering long-term adaptation
Constitutional AI (Bai et al., 2022)	P3	Self-critique generates the revisions at inference (P2)
Feedback-conditioned fine-tuning (Luo et al., 2025)	P3	Critique text retained verbatim in the training context

Table 2: Overview of the three VRL pillars.

Axis	Grounding Signal	Deliberative Feedback	Learning Signal
When	Problem-definition time	Inference time	Training time
Modifies	Task, State, Action, Reward	Outputs or reasoning traces	Model weights, Training data
Persistence	Defines the task	Single episode (extendable via memory)	All future interactions
Methods	Environment specification	Prompting and critique	SFT, PPO, DPO
Challenges	Grounding gap; Compositionality	Feedback Quality; Compute Cost	Signal Compression; Filter Quality

Working Example: the coding agent - The three pillars are complementary, not competing. We emphasize that the pillars classify individual signals, not stages of a pipeline: no method is claimed to traverse a grounding-deliberation-learning cycle, and the example below shows a system in which signals of all three kinds coexist. Consider an agent (Wang et al., 2025d) tasked with resolving a GitHub issue. *Grounding* (Pillar 1) enters first: the natural-language issue description defines the goal, the repository structure specifies the state space, and the test suite implicitly encodes




Figure 3: Grounding of MDP components through natural language. Each subcategory maps verbal feedback to a specific element of the MDP: goal grounding specifies the objective, state grounding represents observations as text, action grounding resolves instructions to executable behavior, and reward code generation compiles language into executable reward functions.

a reward function. At inference time, *deliberative feedback* (Pillar 2) drives the iterative loop: the agent generates a patch, executes the tests, treats traces and error messages as verbal feedback, critiques its own output, and revises. When successful and failed trajectories are later collected and used to fine-tune the model via preference optimization, they become a *learning signal* (Pillar 3) that persistently improves the policy for future tasks. The same agent thus **consumes** verbal feedback at three distinct timescales, each producing a different effect. This example demonstrates the rationale for organizing by temporal scope rather than by feedback source (human, model, tool) or modality (text, scalar, code). Identical verbal feedback can yield fundamentally different effects depending on when it is consumed, and each temporal regime introduces distinct technical challenges, which we will discuss in later sections.

3 Pillar 1: Language as Grounding Signal

In this pillar, we discuss the role of language in providing grounding for each component of the task in which an agent operates. As illustrated in Figure 3, verbal feedback maps to specific MDP elements (Sutton and Barto, 1998): specifying what the agent should optimize (goals - §3.1), what it perceives (states- §3.2), how it can act (actions- §3.3), and what constitutes success (rewards- §3.4) (MacGlashan et al., 2017). This matters because many failures in language-based agents stem not from poor optimization but from poor grounding: feedback that cannot map to executable actions or valid reward conditions will not improve the agent. A common bottleneck across all four subcategories is the grounding gap. More detailed verbal specifications do not necessarily result in improved grounding. The critical factor is whether the mapping from language to MDP components is sufficiently precise for the agent to execute actions.

3.1 Goal Grounding

Goal grounding uses verbal feedback to specify what the agent should achieve, requiring parsing instructions into objects, relations, and target conditions (Chaplot et al., 2018), filtering plans by physical executability (Ahn et al., 2022; Driess et al., 2023), and decomposing goals into verifiable subgoals. Recent work extends this to multi-agent coordination (Zhang et al., 2025d), hierarchical decomposition via offline RL (Hu et al., 2025), and training from language specifications without manual reward functions (Li et al., 2026). The open challenge is compositional generalization: agents must follow goals assembled from familiar primitives in unfamiliar combinations. BabyAI (Chevalier-Boisvert et al., 2018) isolates this in gridworlds, and even state-of-the-art LLMs achieve only about 75% coverage under compositional constraints (Sakai et al., 2025), though compositional policy representations can improve sample efficiency on novel compositions (Cohen et al., 2025).

3.2 State Grounding

Where goal grounding defines what to achieve, state grounding defines what the agent perceives. Verbal state descriptions are most explicit in text-based environments (Yao et al., 2020; Côté et al., 2018), but appear whenever visual or physical states are summarized into language. Shridhar et al. (2020) show that agents trained on verbal state descriptions transfer to embodied 3D settings, demonstrating language as a modality-invariant state representation. Recent work includes closed-loop state feedback for robotic planning (Bhat et al., 2024) and scene-graph-based grounding (Huang et al., 2025a). A key challenge is information preservation, specifically whether the verbal abstraction retains sufficient signal for the policy to operate effectively.

3.3 Action Grounding

Given a goal and a perceived state, the agent must determine how to act. Action grounding resolves verbal instructions to specific skills, tool calls, or motor commands, requiring that instructions map to executable actions and that execution outcomes feedback as language to revise subsequent steps. Huang et al. (2022b) demonstrate this closed loop with Inner Monologue, while Sharma et al. (2022) show that short corrective utterances can amend mistaken steps without disrupting the rest of the

plan. Recent systems have pushed toward tighter integration: multi-modal grounded replanning that corrects ambiguous instructions using visual cues (Kim et al., 2025), constraint-aware visual programming for proactive failure detection (Zhou et al., 2025a), corrective planning that handles physical, logical, and semantic errors across hierarchical control levels (Joublin et al., 2024), and visually grounded task-and-motion planning that combines LLM reasoning with physical feasibility (Zhang et al., 2025c). The central challenge is granularity alignment: the same instruction can refer to a temporally extended skill (“make coffee”), a discrete subgoal (“go to the kitchen”), or a motor command (“close the gripper”), each demanding a different language-to-control interface.

3.4 Reward Code Generation

The preceding subcategories ground objectives, observations, and actions in language. Reward code generation closes the loop by grounding success itself. This subcategory overlaps with Pillar 3 when generated rewards are used for training; we categorize it here because verbal feedback defines the reward function itself, becoming part of the MDP, and Table 1 records the downstream training role explicitly. Under our inclusion criteria, the linguistic signal is consumed at specification time, so reward-code generation satisfies criterion (i) even though the compiled function later emits scalar rewards. The distinguishing feature is compiling verbal descriptions into executable reward functions rather than using language as a direct learning signal. Ma et al. (2024) introduced Eureka, where an LLM generates and iteratively refines reward code based on training summaries, and Xie et al. (2024) extended this to dense reward shaping. This area has seen rapid recent progress: CARD (Sun et al., 2025a) introduces dynamic trajectory-based feedback to refine reward code without per-iteration RL training, PROF (Sun et al., 2025b) extends the paradigm to offline settings through preference optimization, and Zeng et al. (2024) show that reward functions can be generated from video demonstrations. To ensure effectiveness, verbal feedback must be compilable, executable, and accurately reflect the original natural-language intent.




Figure 4: Inference-time verbal feedback mechanisms. Each subcategory exhibits a distinct feedback loop structure: (a) self-critique loops within a single model, (b) externally grounded critique incorporates verified tool or environment output, (c) multi-agent debate distributes critique across parallel model instances and converges to a refined output, (d) experiential memory persists lessons across episodes, and (e) **evaluator feedback in search** explores and prunes multiple candidate paths.

4 Pillar 2: Language as Deliberative Feedback

With the task grounded, we turn to methods that refine reasoning within a single episode without updating model parameters, closely aligning with test-time compute scaling (Snell et al., 2024). As illustrated in Figure 4, five subcategories differ in where critique originates: the model itself (§4.1), external tools (§4.2), peer models (§4.3), past episodes (§4.4), or parallel search paths (§4.5).

4.1 Self-Critique

The simplest deliberative loop asks the model to critique its own output and revise accordingly. Self-Refine (Madaan et al., 2023) formalized this pattern, building on Constitutional AI’s (Bai et al., 2022) demonstration that models can evaluate outputs against written principles. The primary challenge is circularity. Replacing self-generated critique with evaluation from a separately trained critic enhances reasoning (Paul et al., 2024), which indicates that feedback quality is the main bottleneck. Self-correction degrades when the critique source shares the generator’s blind spots (Huang et al., 2024; Kamoi et al., 2024), especially

in smaller models (Olausson et al., 2024). Mitigation strategies include Socratic decomposition into verifiable steps (Shi et al., 2025a) and parallel self-refinement across candidates (Wang et al., 2025c). However, when errors arise from gaps in the model’s knowledge, internal feedback mechanisms cannot identify the missing information.

4.2 Externally Grounded Critique

External grounding addresses this circularity by anchoring verbal feedback in deterministic tool output (Schick et al., 2023). Verified signals, such as execution traces, unit test verdicts, search results, and API responses (Chen et al., 2024; Gou et al., 2024b), facilitate the generation of corrected outputs (Gou et al., 2024a; Yang et al., 2024; Wang et al., 2024b). Additionally, end-to-end reinforcement learning enables models to utilize such feedback more effectively (Gehring et al., 2024). The primary failure mode shifts from blind spots to trust asymmetry: the model receives accurate tool output but may misattribute the root cause of a reported trace. Practical limitations persist, including infrastructure overhead, poor scalability of latency, and the restriction of grounded critique to domains with a verifiable oracle.

4.3 Multi-Agent Critique / Debate

Multi-agent debate distributes verbal critique across parallel model instances, frequently assigning distinct personas to enhance perspective (Hong et al., 2024b). The exchange of critiques can improve reasoning even among identical models (Du et al., 2024), while the introduction of diverse roles further strengthens the quality of discussion (Chan et al., 2023). Communication topology is a critical factor; full-broadcast designs incur substantial token overhead (Choi et al., 2025), motivating the exploration of sparse alternatives (Li et al., 2024). Models may also be trained specifically for debate, enabling them to revise reasoning in response to conflicting opinions (Liu et al., 2026a). Adaptive heterogeneous frameworks have demonstrated improvements in mathematical reasoning (Zhou and Chen, 2025). However, the primary limitation is that distributing critique is only effective when participating models are genuinely diverse. When models share biases, the process yields redundant consensus (Estornell and Liu, 2024; Liang et al., 2024), and current debate methods do not consistently outperform simpler single-agent strategies (Choi et al., 2025).

4.4 Experiential Memory

Previous approaches regenerate feedback from scratch in each episode. In contrast, experiential memory distills reusable lessons into persistent storage, categorized as episodic (raw logs), semantic (distilled rules), and procedural (reusable skills) (Zhang et al., 2025e). **Although stored lessons persist across episodes, we classify memory by the moment of consumption: a retrieved lesson acts as deliberative feedback within the current episode, and Table 1 records the boundary with long-term adaptation.** Reflexion (Shinn et al., 2023) stores raw verbal feedback, ExpeL (Zhao et al., 2024) extracts insights from paired successes and failures, and Voyager (Wang et al., 2023a) accumulates executable skills. Recent surveys have formalized these mechanisms (Du, 2026), while MemBench evaluates memory quality through downstream performance (Tan et al., 2025). However, persisting feedback can also perpetuate errors: poor-quality memories propagate mistakes, expanding memory stores require complex retrieval (Park et al., 2023), outdated knowledge can mislead in non-stationary environments (Zhang et al., 2025e), and adversarial injection may induce cross-session drift (Dong et al., 2026).

4.5 Evaluator Feedback in Search

The inclusion criterion for this subcategory is the verbal evaluator signal consumed during search, not the search procedure itself; decoding strategies that involve no linguistic feedback fall outside VRL. While previous approaches refine a single trajectory, **these methods** explore multiple candidate paths and employ verbal feedback to determine which paths to expand, prune, or abandon (Wang et al., 2025a). The Tree of Thoughts framework (Yao et al., 2023) samples several intermediate thoughts and evaluates them through self-assessment, thereby extending chain-of-thought prompting (Wei et al., 2022) with a tree-based search mechanism. Reasoning via Planning (Hao et al., 2023) and Graph of Thoughts (Besta et al., 2024) further generalize the underlying topology, permitting non-linear dependencies such as aggregation and refinement. A key advantage of these methods is that verbal feedback can be applied at every thought level, which facilitates earlier identification of unproductive paths. However, exploring multiple states necessitates numerous large language model (LLM) calls, rendering these methods

considerably more expensive than single-pass generation (Xie et al., 2023).

5 Pillar 3: Language as Learning Signal

The method above leave the model unchanged; we now examine how verbal feedback can persistently reshape the policy through training. As illustrated in Figure 5, the four subcategories form a compression spectrum: from feedback-conditioned modeling (§5.1), which preserves full critiques, through self-improvement (§5.2) and process supervision (§5.3), to preference shaping (§5.4), which reduces verbal judgement to a single scalar.

5.1 Feedback-Conditioned Modeling

At the upper end of the compression spectrum, feedback-conditioned modeling maintains verbal critiques in their entirety by conditioning the model directly on natural-language feedback during training. The training signal consists of a triple (x, v, a^*) : the input, a critique of an initial response, and the revised output. A central question is whether to retain the critique within the training context. Early studies omitted the critique, using only the refined output as the supervision target (Scheurer et al., 2023). Subsequent work, such as ALT (Lloret et al., 2024), demonstrated that retaining feedback as conditioning context enables models to learn explicit mappings from critique to correction. Luo et al. (2025) further formalized this approach by treating free-form feedback as a primary conditioning signal with online bootstrapping. While preserving the full linguistic content of the critique allows these methods to retain the richest signal, this same receptiveness introduces a risk: models may learn to follow critiques indiscriminately, even when they are incorrect (Sharma et al., 2024).

5.2 Self-Improvement

Self-improvement approaches treat language as filtered text rather than as a conditioning context. The central mechanism is generate-then-filter: the model generates multiple candidate trajectories, verbal feedback signals determine which candidates are retained, and these selected outputs become supervised fine-tuning data for subsequent iterations (Singh et al., 2023). The filtering criterion directly shapes the model’s learning. For example, in STaR (Zelikman et al., 2022), correctness of the final answer is used as the filter, while in Constitutional AI (Bai et al., 2022), adherence




Figure 5: Training-time verbal feedback methods organized by degree of linguistic compression. Methods at the top preserve full natural-language feedback during training, while methods at the bottom compress it into scalar signals.

to safety principles serves this role. Explicit labeling is not always required; majority voting across diverse samples can identify likely correct traces (Huang et al., 2022a), and the model itself can act as a judge across iterations (Yuan et al., 2024). This self-contained loop can facilitate compounding improvements over multiple cycles. However, filter quality remains a critical bottleneck: when the same model functions as both generator and judge, systematic biases may be amplified (Shumailov et al., 2023).

5.3 Process Supervision

Process supervision reduces verbal feedback to scalar scores at the step level. Annotators assess each intermediate reasoning step, and these evaluations are used to train a process reward model; the verbal content of the judgments is omitted. Despite this reduction, step-level supervision substantially outperforms outcome-level supervision (Lightman et al., 2024; Uesato et al., 2022), as localized feedback enables more precise credit assignment. The high cost of step-level annotation has motivated efforts to automate this process (Wang et al., 2024a; Luo et al., 2024). Generative process reward models that incorporate reasoning prior to scoring retain more verbal information while maintaining step-level granularity (Zhao et al., 2026; Liu et al., 2025b). Although step-level granularity offers significant advantages, it remains costly and is challenging to define in domains where intermediate correctness is ambiguous.

5.4 Preference Shaping

Preference shaping represents the maximal compression of verbal feedback, reducing a comparative judgment between two responses to a single scalar that guides policy updates. Building on the foundation of learning from verbal preferences

(Christiano et al., 2017), InstructGPT (Ouyang et al., 2022) demonstrated large-scale implementation by training a reward model on pairwise rankings and optimizing the policy online. Direct Preference Optimization (DPO) (Rafailov et al., 2023) further simplified this process by eliminating the explicit reward model and training directly on preference pairs, using the comparison itself as the loss signal. Subsequent research has investigated additional alternatives (Ethayarajh et al., 2024; Hong et al., 2024a; Meng et al., 2024; Azar et al., 2024). Notably, Lee et al. (2023) replaced human annotators with AI judges, thereby enhancing the scalability of preference collection. The online and offline variants differ in practice: online optimization against a learned reward model can query feedback on fresh policy samples, while offline objectives such as DPO inherit the coverage of a fixed preference set. Although scalar compression reduces the explanatory richness of verbal feedback, the scalability enabled by llm generated judgments increasingly justifies this tradeoff. Under our inclusion criteria (Section 1), preference shaping marks the outer boundary of VRL: once the judgment is compressed to a preference bit before any learning, the linguistic structure is no longer consumed, so vanilla RLHF and DPO sit outside the paradigm. We retain them here as the endpoint of the compression spectrum because they clarify, by contrast, what the methods above preserve.

6 Insights, Discussion and Future Prospects

In practice, the three pillars operate together, as the coding-agent example in §2 illustrates. Five cross-cutting themes span them: algorithmic interactions and practitioner guidance (§6.1), feedback-model quality (§6.2), tool-interface design (§6.3), adversarial robustness (§6.4), and the absence of formal guarantees (§6.5).

6.1 Algorithmic Interactions and Practitioner Guidance

Different feedback forms pair with different optimization families, and the pairing carries assumptions a practitioner must check. Executable reward functions generated from language pair with online policy-gradient training; pairwise judgments pair with offline preference objectives such as DPO; filtered or feedback-conditioned text pairs with supervised fine-tuning; and step-level verbal judgments

pair with process reward models (PRMs), which enable finer-grained credit assignment (Lightman et al., 2024). Each pairing has consequences inside the decision problem: rewards specified through language inherit objective mis-specification and reward-hacking risks; states represented through language raise information-loss and aliasing questions; and actions generated under LLM priors change exploration. Appendix B distills these considerations into a decision table keyed to properties a practitioner can assess before committing to a mechanism: verifiability of task outcomes, feedback latency and cost, and risk exposure such as reward hacking, sycophancy, and feedback poisoning.

6.2 Dedicated Feedback Models

Verbal feedback quality remains the key performance factor, but current practice typically reuses the same LLM for both generation and critique. Dedicated critics consistently outperform this approach: Shepherd (Wang et al., 2023b) provides more actionable feedback than untuned baselines, CriticGPT (McAleese et al., 2024) detects code bugs more reliably, and CritiqueLLM (Ke et al., 2024) scales critique data via multi-path prompting. Co-training critics with reasoners yields further improvements (Liu et al., 2025a), while generative process reward models that reason before scoring preserve more verbal signal than scalar PRMs (Khalifa et al., 2025). CriticBench (Lin et al., 2024) offers the first systematic benchmark for critique quality. A broader evaluation practice is emerging around reward-model benchmarks (Lambert et al., 2024) and step-level process benchmarks (Zheng et al., 2024); we recommend reporting feedback quality and feedback utilization separately, since models often fail to incorporate even accurate feedback (Jiang et al., 2025). These results indicate a natural next step: just as the field invested in instruction-tuned models (Ouyang et al., 2022), it should now develop feedback-tuned models whose objectives address error localization, actionability, and calibration. Recent advances in span-level feedback gradients (Wang et al., 2025b) and multi-turn didactic interactions (Klissarov et al., 2026) suggest such models can serve as differentiable training signals across each pillars.

6.3 Designing Tools for Verbal Feedback

Developer-oriented tools often yield outputs that are terse or conflate error types. Yang et al.

(2024) demonstrated that LLM-optimized interfaces enhance agent performance, while Bandlamudi et al. (2025) showed ambiguous outputs cause many agent failures across 750 API calls. Tool-interactive critiquing requires interpretable output (Gou et al., 2024a), and executable code actions benefit agents by providing structured feedback (Wang et al., 2024b). RL on execution feedback (Gehring et al., 2024) presumes adequate tool-output signal, while agent harness frameworks (Ning et al., 2026) set tool quality as the limit of iterative refinement. However, a metric to measure this limit is lacking: tool-output compliance should quantify agent consumability. Tool APIs must expose structured metadata distinguishing task-level errors from infrastructure failures to enable agents to apply suitable recovery strategies.

6.4 Adversarial Verbal Feedback

Because VRL agents inherently act on feedback, adversarial feedback constitutes policy manipulation, introducing vulnerabilities beyond standard prompt injection. Zhan et al. (2025) demonstrated that adaptive attacks circumvent all eight tested defenses, while tool-level injection attains a 96.7% success rate on GPT-4o (Shi et al., 2025b). Hidden instructions in API responses covertly manipulate agents (Zhan et al., 2024), adversarial memory injection induces cross-session drift (Dong et al., 2026), and even benign incorrect feedback leads to sycophancy (Sharma et al., 2024). Defenses are developing: instruction hierarchies (Wallace et al., 2024), adversarial fine-tuning (Chen et al., 2025), and system-level design patterns (Beurer-Kellner et al., 2025) each mitigate narrow attack surfaces. Comprehensive solutions will require feedback-provenance mechanisms to verify source integrity and assign trust, as well as adversarial feedback benchmarks to evaluate robustness as the primary metric.

6.5 Theoretical Foundations

Verbal RL lacks formal guarantees; it remains unclear when verbal feedback enhances sample efficiency or policy quality. Three promising directions emerge. First, if verbal critics serve as noisy oracles with bounded error, PAC-learning frameworks (Strehl et al., 2009) can bound VRL sample complexity, and critic benchmarks (Lin et al., 2024) allow empirical estimation of error rates. Second, modeling feedback as partial observation links VRL to POMDP instruction following

(Luketina et al., 2019), where belief-state planning yields convergence guarantees; McCallum et al. (2023) implement this with feedback-conditioned decision transformers. Third, rate-distortion theory may quantify signal retention from critique to scalar, clarifying when preference shaping aligns with feedback-conditioned methods. Empirical studies support these directions: language models learn from verbal feedback without scalar rewards (Luo et al., 2025), text feedback broadens RL capabilities (Song et al., 2026), and recent surveys link RLHF to language-conditioned RL theory (Kaufmann et al., 2023). Formalizing results in these areas would clarify when the cost of rich verbal feedback is warranted over scalar alternatives. **We state these as open problems rather than established results. Concretely, the field’s research agenda includes: the sample-efficiency conditions under which verbal feedback beats scalar reward; optimality and convergence when rewards are specified or shaped through language; temporal credit assignment from free-form critiques; provenance and robustness of feedback channels; and standardized, separate evaluation of feedback quality and feedback utilization. Each item is currently blocked by the absence of a formal account of what a verbal signal contributes beyond its scalarization.**

7 Conclusion

In this paper, we present the first comprehensive survey of verbal reinforcement learning, organized around three pillars based on when language enters an agent’s lifecycle: grounding signal, deliberative feedback, and learning signal. Together, these pillars reveal a common pattern: verbal feedback is rapidly becoming the primary medium for defining, updating, and improving agents.

Looking ahead, we anticipate that verbal feedback will become a first-class design primitive in agent architectures. The boundaries between pillars will increasingly blur as agents consume verbal feedback in unified loops that span grounding, deliberation, and learning within a single trajectory. The bottleneck will shift from *generating* feedback to *verifying* it, making feedback provenance, quality benchmarks, and adversarial robustness first-class infrastructure requirements. Finally, tool interfaces will need to be redesigned around agent consumability rather than human readability, as the gap between tool output quality and agent performance becomes the practical ceiling for iterative

refinement.

Limitations

We aim to provide broad coverage of methods that use verbal feedback within the agent’s lifecycle. Within each category, we focus on representative papers rather than attempting an exhaustive enumeration. Additionally, while our taxonomy offers a unifying framework, it may oversimplify methods that span multiple roles of language. Several approaches do not fit neatly into a single branch, since verbal feedback can simultaneously serve as a grounding signal, a deliberative mechanism, and a learning signal. Our taxonomy should therefore be understood as an organizing lens rather than a strict partition of the literature. Finally, our survey focuses primarily on methods in which verbal feedback is explicit and central to the agent’s process, and therefore does not fully cover adjacent work where language plays a more auxiliary role. **Language-based feedback itself carries failure modes that this survey can name but not yet bound: reward-specification failures when language compiles to the wrong objective, unreliable interpretation of feedback by the consuming model, weak grounding between verbal descriptions and environment state, and the transfer of language-guided policies to real-world settings. Societal impact follows the same channels: when feedback is incorrect, biased, adversarial, or poorly grounded, harm falls on users of high-stakes deployments such as coding assistants, robotics, clinical decision support, and education, and on the people affected by those systems’ outputs. Sections 6.4 and 6.5 discuss mitigation directions.** AI writing tools were used for language editing and proofreading only; all research contributions are the authors’ own.

References

- Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altschmidt, Sam Altman, and Shyamal Anadkat. 2023. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*.
- Michael Ahn, Anthony Brohan, Noah Brown, Yevgen Chebotar, Omar Cortes, Byron David, Chelsea Finn, Chuyuan Fu, Keerthana Gopalakrishnan, and Karol Hausman. 2022. Do as i can, not as i say: Grounding language in robotic affordances. *arXiv preprint arXiv:2204.01691*.
- Bakhtawar Ahtisham, Kirk Vanacore, Jinsook Lee, Zhuqian Zhou, Doug Pietrzak, and Rene F Kizil-

847	cec. 2026. Ai annotation orchestration: Evaluating	Chi-Min Chan, Weize Chen, Yusheng Su, Jianxuan Yu,	904
848	llm verifiers to improve the quality of llm annota-	Wei Xue, Shanghang Zhang, Jie Fu, and Zhiyuan	905
849	tions in learning analytics. In <i>Proceedings of the</i>	Liu. 2023. Chateval: Towards better llm-based	906
850	<i>LAK26: 16th International Learning Analytics and</i>	evaluators through multi-agent debate . <i>Preprint</i> ,	907
851	<i>Knowledge Conference</i> , pages 447–456.	arXiv:2308.07201.	908
852	Dario Amodei, Chris Olah, Jacob Steinhardt, Paul	Devendra Singh Chaplot, Kanthashree Mysore Sathyen-	909
853	Christiano, John Schulman, and Dan Mané. 2016.	dra, Rama Kumar Pasumarthi, Dheeraj Rajagopal,	910
854	Concrete problems in ai safety. <i>arXiv preprint</i>	and Ruslan Salakhutdinov. 2018. Gated-attention	911
855	<i>arXiv:1606.06565</i> .	architectures for task-oriented language grounding.	912
856	Mohammad Gheshlaghi Azar, Zhaohan Daniel Guo, Bi-	In <i>Proceedings of the AAAI Conference on Artificial</i>	913
857	lial Piot, Remi Munos, Mark Rowland, Michal Valko,	<i>Intelligence</i> , volume 32.	914
858	and Daniele Calandriello. 2024. A general theoret-	Sizhe Chen, Arman Zharmagambetov, Saeed Mahlouji-	915
859	ical paradigm to understand learning from human	far, Kamalika Chaudhuri, David Wagner, and Chuan	916
860	preferences. In <i>International Conference on Arti-</i>	Guo. 2025. Secalign: Defending against prompt	917
861	<i>ficial Intelligence and Statistics</i> , pages 4447–4455.	injection with preference optimization . <i>Preprint</i> ,	918
862	PMLR.	arXiv:2410.05451.	919
863	Yuntao Bai, Saurav Kadavath, Sandipan Kundu,	Ting Chen, Simon Kornblith, Mohammad Norouzi, and	920
864	Amanda Askell, Jackson Kernion, Andy Jones, Anna	Geoffrey Hinton. 2020. A simple framework for	921
865	Chen, Anna Goldie, Azalia Mirhoseini, Cameron	contrastive learning of visual representations . In <i>In-</i>	922
866	McKinnon, Carol Chen, Catherine Olsson, Christo-	<i>ternational conference on machine learning</i> , pages	923
867	pher Olah, Danny Hernandez, Dawn Drain, Deep	1597–1607. Pmlr.	924
868	Ganguli, Dustin Li, Eli Tran-Johnson, Ethan Perez,	Xinyun Chen, Maxwell Lin, Nathanael Schärli, and	925
869	and 32 others. 2022. Constitutional AI: Harmlessness	Denny Zhou. 2024. Teaching large language models	926
870	from AI feedback. <i>arXiv preprint arXiv:2212.08073</i> .	to self-debug. In <i>International Conference on Learn-</i>	927
871	Jayachandu Bandlamudi, Ritwik Chaudhuri, Neelamad-	<i>ing Representations</i> , volume 2024, pages 8746–8825.	928
872	hav Gantayat, Sambit Ghosh, Kushal Mukherjee, Pre-	Maxime Chevalier-Boisvert, Dzmitry Bahdanau, Salem	929
873	rna Agarwal, Renuka Sindhgatta, and Sameep Mehta.	Lahlou, Lucas Willems, Chitwan Saharia, Thien Huu	930
874	2025. A framework for testing and adapting rest apis	Nguyen, and Yoshua Bengio. 2018. Babyai: A plat-	931
875	as llm tools. <i>arXiv preprint arXiv:2504.15546</i> .	form to study the sample efficiency of grounded lan-	932
876	Maciej Besta, Nils Blach, Ales Kubicek, Robert Gersten-	guage learning. <i>arXiv preprint arXiv:1810.08272</i> .	933
877	berger, Michal Podstawski, Lukas Gianinazzi, Joanna	Hyeong Kyu Choi, Xiaojin Zhu, and Yixuan Li. 2025.	934
878	Gajda, Tomasz Lehmann, Hubert Niewiadomski, and	Debate or vote: Which yields better decisions in	935
879	Piotr Nyczyk. 2024. Graph of thoughts: Solving	multi-agent large language models? <i>Advances in</i>	936
880	elaborate problems with large language models. In	<i>Neural Information Processing Systems</i> .	937
881	<i>Proceedings of the AAAI conference on artificial in-</i>	Paul F Christiano, Jan Leike, Tom Brown, Miljan Mar-	938
882	<i>telligence</i> , volume 38, pages 17682–17690.	tic, Shane Legg, and Dario Amodei. 2017. Deep	939
883	Luca Beurer-Kellner, Beat Buesser, Ana-Maria Crețu,	reinforcement learning from human preferences . <i>Ad-</i>	940
884	Edoardo Debenedetti, Daniel Dobos, Daniel Fabian,	<i>vances in Neural Information Processing Systems</i> ,	941
885	Marc Fischer, David Froelicher, Kathrin Grosse, and	30.	942
886	Daniel Naeff. 2025. Design patterns for securing	Vanya Cohen, Geraud Nangue Tasse, Nakul Gopalan,	943
887	llm agents against prompt injections. <i>arXiv preprint</i>	Steven James, Matthew Gombolay, Ray Mooney, and	944
888	<i>arXiv:2506.08837</i> .	Benjamin Rosman. 2025. Compositional instruction	945
889	Vineet Bhat, Ali Umut Kaypak, Prashanth Krish-	following with language models and reinforcement	946
890	namurthy, Ramesh Karri, and Farshad Khorrami.	learning. <i>arXiv preprint arXiv:2501.12539</i> .	947
891	2024. Grounding llms for robot task planning	Marc-Alexandre Côté, Akos Kádár, Xingdi Yuan, Ben	948
892	using closed-loop state feedback. <i>arXiv preprint</i>	Kybartas, Tavian Barnes, Emery Fine, James Moore,	949
893	<i>arXiv:2402.08546</i> .	Matthew Hausknecht, Layla El Asri, and Mahmoud	950
894	Andres Bran, Sam Cox, Oliver Schilter, Carlo Baldas-	Adada. 2018. Textworld: A learning environment for	951
895	sari, Andrew D White, and Philippe Schwaller. 2024.	text-based games. In <i>Workshop on Computer Games</i> ,	952
896	Augmenting large language models with chemistry	pages 41–75. Springer.	953
897	tools. <i>Nature machine intelligence</i> , 6(5):525–535.	Jonas Degraeve, Federico Felici, Jonas Buchli, Michael	954
898	Tom Brown, Benjamin Mann, Nick Ryder, Melanie Sub-	Neunert, Brendan Tracey, Francesco Carpanese,	955
899	biah, Jared D Kaplan, Prafulla Dhariwal, Arvind Nee-	Timo Ewalds, Roland Hafner, Abbas Abdolmaleki,	956
900	lakantan, Pranav Shyam, Girish Sastry, and Amanda	and Diego de Las Casas. 2022. Magnetic control of	957
901	Askell. 2020. Language models are few-shot learners.	tokamak plasmas through deep reinforcement learn-	958
902	<i>Advances in neural information processing systems</i> ,	ing. <i>Nature</i> , 602(7897):414–419.	959
903	33:1877–1901.		

960	Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. Bert: Pre-training of deep bidirectional transformers for language understanding . <i>Preprint</i> , arXiv:1810.04805.	1015
961		1016
962		1017
963		1018
964	Shen Dong, Shaochen Xu, Pengfei He, Yige Li, Jiliang Tang, Tianming Liu, Hui Liu, and Zhen Xiang. 2026. Memory injection attacks on llm agents via query-only interaction. <i>Advances in Neural Information Processing Systems</i> , 38:46697–46731.	1019
965		
966		1020
967		1021
968		1022
969	Danny Driess, Fei Xia, Mehdi SM Sajjadi, Corey Lynch, Aakanksha Chowdhery, Brian Ichter, Ayzaan Wahid, Jonathan Tompson, Quan Vuong, and Tianhe Yu. 2023. Palm-e: An embodied multimodal language model. <i>arXiv preprint arXiv:2303.03378</i> .	1023
970		1024
971		
972		1025
973		1026
974	Pengfei Du. 2026. Memory for autonomous llm agents: Mechanisms, evaluation, and emerging frontiers. <i>arXiv preprint arXiv:2603.07670</i> .	1027
975		1028
976		1029
977	Yilun Du, Shuang Li, Antonio Torralba, Joshua B Tenenbaum, and Igor Mordatch. 2024. Improving factuality and reasoning in language models through multi-agent debate. In <i>International Conference on Machine Learning</i> .	1030
978		
979		1031
980		1032
981		1033
982	Andrew Estornell and Yang Liu. 2024. Multi-llm debate: Framework, principals, and interventions. <i>Advances in Neural Information Processing Systems</i> , 37:28938–28964.	1034
983		
984		1035
985		1036
986	Kawin Ethayarajh, Winnie Xu, Niklas Muennighoff, Dan Jurafsky, and Douwe Kiela. 2024. Kto: Model alignment as prospect theoretic optimization. <i>arXiv preprint arXiv:2402.01306</i> .	1037
987		1038
988		1039
989		
990	Jinyuan Fang, Yanwen Peng, Xi Zhang, Yingxu Wang, Xinhao Yi, Guibin Zhang, and 1 others. 2025. A comprehensive survey of self-evolving AI agents: A new paradigm bridging foundation models and lifelong agentic systems. <i>arXiv preprint arXiv:2508.07407</i> .	1040
991		1041
992		1042
993		1043
994		1044
995	Xidong Feng, Bo Liu, Yan Song, Haotian Fu, Ziyu Wan, Girish A. Koushik, and 1 others. 2024. Natural Language Reinforcement Learning. <i>arXiv preprint arXiv:2411.14251</i> .	1045
996		1046
997		
998		1047
999	Patrick Fernandes, Aman Madaan, Emmy Liu, António Farinhas, Pedro Henrique Martins, Amanda Bertsch, and 1 others. 2023. Bridging the gap: A survey on integrating (human) feedback for natural language generation. <i>Transactions of the Association for Computational Linguistics</i> , 11. ArXiv:2305.00955.	1048
1000		1049
1001		1050
1002		1051
1003		
1004		1052
1005	Huan-ang Gao, Jiayi Geng, Wenyue Hua, Mengkang Hu, Xinzhe Juan, Hongzhang Liu, and 1 others. 2025. A survey of self-evolving agents: What, when, how, and where to evolve on the path to artificial super intelligence. <i>arXiv preprint arXiv:2507.21046</i> .	1053
1006		1054
1007		1055
1008		1056
1009		
1010	Jonas Gehring, Kunhao Zheng, Jade Copet, Vegard Mella, Quentin Carbonneaux, Taco Cohen, and Gabriel Synnaeve. 2024. Rlef: Grounding code llms in execution feedback with reinforcement learning. <i>arXiv preprint arXiv:2410.02089</i> .	1057
1011		1058
1012		1059
1013		1060
1014		
	Zhibin Gou, Zhihong Shao, Yeyun Gong, Yelong Shen, Yujiu Yang, Nan Duan, and Weizhu Chen. 2024a. CRITIC: Large language models can self-correct with tool-interactive critiquing . In <i>International Conference on Learning Representations</i> .	1061
		1062
		1063
		1064
		1065
		1066
	Zhibin Gou, Zhihong Shao, Yeyun Gong, Yelong Shen, Yujiu Yang, Minlie Huang, Nan Duan, and Weizhu Chen. 2024b. Tora: A tool-integrated reasoning agent for mathematical problem solving . <i>Preprint</i> , arXiv:2309.17452.	1067
		1068
		1069
		1070
	Shibo Hao, Yi Gu, Haodi Ma, Joshua Hong, Zhen Wang, Daisy Wang, and Zhiting Hu. 2023. Reasoning with language model is planning with world model. In <i>Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing</i> , pages 8154–8173.	
	Ali Hatamizadeh, Shrimai Prabhumoye, Igor Gitman, Ximing Lu, Seungju Han, Wei Ping, Yejin Choi, and Jan Kautz. 2026. igrpo: Self-feedback-driven llm reasoning . <i>arXiv preprint arXiv:2602.09000</i> .	
	Jiwoo Hong, Noah Lee, and James Thorne. 2024a. Orpo: Monolithic preference optimization without reference model. In <i>Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing</i> , pages 11170–11189.	
	Sirui Hong, Mingchen Zhuge, Jonathan Chen, Xiawu Zheng, Yuheng Cheng, Jinlin Wang, Ceyao Zhang, Steven Yau, Zijuan Lin, and Liyang Zhou. 2024b. Metagpt: Meta programming for a multi-agent collaborative framework. In <i>International Conference on Learning Representations</i> , volume 2024, pages 23247–23275.	
	Zican Hu, Wei Liu, Xiaoye Qu, Xiangyu Yue, Chunlin Chen, Zhi Wang, and Yu Cheng. 2025. Divide and conquer: Grounding llms as efficient decision-making agents via offline hierarchical reinforcement learning. <i>arXiv preprint arXiv:2505.19761</i> .	
	Jiani Huang, Amish Sethi, Matthew Kuo, Mayank Keoliya, Neelay Velingker, JungHo Jung, Ser-Nam Lim, Ziyang Li, and Mayur Naik. 2025a. Esca: Contextualizing embodied agents via scene-graph generation . <i>Preprint</i> , arXiv:2510.15963.	
	Jiaxin Huang, Shixiang Shane Gu, Le Hou, Yuexin Wu, Xuezhi Wang, Hongkun Yu, and Jiawei Han. 2022a. Large language models can self-improve. <i>arXiv preprint arXiv:2210.11610</i> .	
	Jie Huang, Xinyun Chen, Swaroop Mishra, Huaixiu Steven Zheng, Adams Yu, Xinying Song, and Denny Zhou. 2024. Large language models cannot self-correct reasoning yet. In <i>International conference on learning representations</i> , volume 2024, pages 32808–32824.	
	Wenlong Huang, Fei Xia, Ted Xiao, Harris Chan, Jacky Liang, Pete Florence, Andy Zeng, Jonathan Tompson, Igor Mordatch, and Yevgen Chebotar. 2022b. Inner monologue: Embodied reasoning through	

1071	planning with language models. <i>arXiv preprint</i>	Nathan Lambert, Valentina Pyatkin, Jacob Morrison,	1125
1072	<i>arXiv:2207.05608</i> .	LJ Miranda, Bill Yuchen Lin, Khyathi Chandu,	1126
1073	Yue Huang, Yanyuan Chen, Dexuan Xu, Weihua	and 1 others. 2024. RewardBench: Evaluating re-	1127
1074	Yue, Huamin Zhang, Meikang Qiu, and Yu Huang.	ward models for language modeling. <i>arXiv preprint</i>	1128
1075	2025b. Medreflect: Teaching medical llms to self-	<i>arXiv:2403.13787</i> .	1129
1076	improve via reflective correction. <i>arXiv preprint</i>		
1077	<i>arXiv:2510.03687</i> .	Harrison Lee, Samrat Phatale, Hassan Mansoor, Thomas	1130
1078	Dongwei Jiang, Alvin Zhang, Andrew Wang, Nicholas	Mesnard, Johan Ferret, Kellie Lu, Colton Bishop,	1131
1079	Andrews, and Daniel Khashabi. 2025. Feedback	Ethan Hall, Victor Carbune, and Abhinav Rastogi.	1132
1080	friction: LLMs struggle to fully incorporate external	2023. Rlaif vs. rlhf: Scaling reinforcement learn-	1133
1081	feedback. <i>arXiv preprint arXiv:2506.11930</i> .	ing from human feedback with ai feedback. <i>arXiv</i>	1134
1082		<i>preprint arXiv:2309.00267</i> .	1135
1083	Carlos E. Jimenez, John Yang, Alexander Wettig,	Andrew Li, Toryn Klassen, Andrew Wang, Parand	1136
1084	Shunyu Yao, Kexin Pei, Ofir Press, and Karthik	A Alamdari, and Sheila McIlraith. 2026. Ground-	1137
1085	Narasimhan. 2024. Swe-bench: Can language mod-	compose-reinforce: Grounding language in agentic	1138
1086	els resolve real-world github issues? <i>Preprint</i> ,	behaviours using limited data. <i>Advances in Neural</i>	1139
1087	<i>arXiv:2310.06770</i> .	<i>Information Processing Systems</i> , 38:160838–160874.	1140
1088	Frank Joublin, Antonello Ceravola, Pavel Smirnov, Fe-	Yunxuan Li, Yibing Du, Jiageng Zhang, Le Hou, Peter	1141
1089	lix Ocker, Joerg Deigmoeller, Anna Belardinelli,	Grabowski, Yeqing Li, and Eugene Ie. 2024. Improv-	1142
1090	Chao Wang, Stephan Hasler, Daniel Tanneberg, and	ing multi-agent debate with sparse communication	1143
1091	Michael Gienger. 2024. Copal: corrective planning	topology. In <i>Findings of the Association for Com-</i>	1144
1092	of robot actions with large language models. In <i>2024</i>	<i>putational Linguistics: EMNLP 2024</i> , pages 7281–	1145
1093	<i>ieee international conference on robotics and au-</i>	7294.	1146
1094	<i>tomation (ICRA)</i> , pages 8664–8670. IEEE.	Tian Liang, Zhiwei He, Wenxiang Jiao, Xing Wang,	1147
1095	Ryo Kamoi, Yusen Zhang, Nan Zhang, Jiawei Han,	Yan Wang, Rui Wang, Yujiu Yang, Shuming Shi, and	1148
1096	and Rui Zhang. 2024. When can llms actually cor-	Zhaopeng Tu. 2024. Encouraging divergent thinking	1149
1097	rect their own mistakes? a critical survey of self-	in large language models through multi-agent debate .	1150
1098	correction of llms. <i>Transactions of the Association</i>	<i>Preprint</i> , arXiv:2305.19118.	1151
1099	<i>for Computational Linguistics</i> , 12:1417–1440.	Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harri-	1152
1100	Timo Kaufmann, Paul Weng, Viktor Bengs, and Eyke	son Edwards, Bowen Baker, Teddy Lee, Jan Leike,	1153
1101	Hüllermeier. 2023. A survey of reinforcement	John Schulman, Ilya Sutskever, and Karl Cobbe.	1154
1102	learning from human feedback. <i>arXiv preprint</i>	2024. Let’s verify step by step . In <i>International</i>	1155
1103	<i>arXiv:2312.14925</i> .	<i>Conference on Learning Representations</i> , volume	1156
1104	Pei Ke, Bosi Wen, Andrew Feng, Xiao Liu, Xuanyu Lei,	2024, pages 39578–39601.	1157
1105	Jiale Cheng, Shengyuan Wang, Aohan Zeng, Yux-	Zicheng Lin, Zhibin Gou, Tian Liang, Ruilin Luo,	1158
1106	iao Dong, and Hongning Wang. 2024. Critiquellm:	Haowei Liu, and Yujiu Yang. 2024. Criticbench:	1159
1107	Towards an informative critique generation model	Benchmarking llms for critique-correct reasoning. In	1160
1108	for evaluation of large language model generation.	<i>Findings of the Association for Computational Lin-</i>	1161
1109	In <i>Proceedings of the 62nd Annual Meeting of the</i>	<i>guistics: ACL 2024</i> , pages 1552–1587.	1162
1110	<i>Association for Computational Linguistics (Volume</i>	Chenxi Liu, Yanshuo Chen, Ruibo Chen, Tianyi Xiong,	1163
1111	<i>1: Long Papers)</i> , pages 13034–13054.	Tong Zheng, and Heng Huang. 2026a. Learning from	1164
1112	Muhammad Khalifa, Rishabh Agarwal, Lajanugen Lo-	self-debate: Preparing reasoning models for multi-	1165
1113	geswaran, Jaekyeom Kim, Hao Peng, Moontae Lee,	agent debate . <i>Preprint</i> , arXiv:2601.22297.	1166
1114	Honglak Lee, and Lu Wang. 2025. Process reward	Qihao Liu, Luoxin Ye, Wufei Ma, Yu-Cheng Chou, and	1167
1115	models that think . <i>Preprint</i> , arXiv:2504.16828.	Alan Yuille. 2025a. Generative adversarial reasoner:	1168
1116	Taewoong Kim, Byeonghwi Kim, and Jonghyun Choi.	Enhancing llm reasoning with adversarial reinforce-	1169
1117	2025. Multi-modal grounded planning and efficient	ment learning. <i>arXiv preprint arXiv:2512.16917</i> .	1170
1118	replanning for learning embodied agents with a few	ShaoZhen Liu, Xinting Huang, Houwen Peng, Xin	1171
1119	examples. In <i>Proceedings of the AAAI Conference on</i>	Chen, Xinyang Song, Qi Li, and Zhenan Sun. 2026b.	1172
1120	<i>Artificial Intelligence</i> , volume 39, pages 4329–4337.	Dual-phase llm reasoning: Self-evolved mathemati-	1173
1121	Martin Klissarov, Jonathan Cook, Diego Antognini, Hao	cal frameworks. <i>arXiv preprint arXiv:2601.05616</i> .	1174
1122	Sun, Jingling Li, Natasha Jaques, Claudiu Musat, and	Xiaoqian Liu, Ke Wang, Yuchuan Wu, Fei Huang, Yong-	1175
1123	Edward Grefenstette. 2026. Improving interactive	bin Li, Junge Zhang, and Jianbin Jiao. 2025b. Agen-	1176
1124	in-context learning from natural language feedback.	tic reinforcement learning with implicit step rewards .	1177
	<i>arXiv preprint arXiv:2602.16066</i> .	<i>Preprint</i> , arXiv:2509.19199.	1178

1179	Saiic Abadal Lloret, Shehzaad Dhuliawala, Keerthiram	Yu Meng, Mengzhou Xia, and Danqi Chen. 2024.	1235
1180	Murugesan, and Mrinmaya Sachan. 2024. Towards	SimpO: Simple preference optimization with a	1236
1181	aligning language models with textual feedback. In	reference-free reward. <i>Advances in Neural Infor-</i>	1237
1182	<i>Proceedings of the 2024 Conference on Empirical</i>	<i>mation Processing Systems</i> , 37:124198–124235.	1238
1183	<i>Methods in Natural Language Processing</i> , pages		
1184	20240–20266.		
1185	Chris Lu, Cong Lu, Robert Tjarko Lange, Jakob Foer-	Azalia Mirhoseini, Anna Goldie, Mustafa Yazgan,	1239
1186	ster, Jeff Clune, and David Ha. 2024. The ai scientist:	Joe Wenjie Jiang, Ebrahim Songhori, Shen Wang,	1240
1187	Towards fully automated open-ended scientific dis-	Young-Joon Lee, Eric Johnson, Omkar Pathak, and	1241
1188	covery. <i>arXiv preprint arXiv:2408.06292</i> .	Azade Nova. 2021. A graph placement methodology	1242
1189	Jelena Luketina, Nantas Nardelli, Gregory Farquhar,	for fast chip design. <i>Nature</i> , 594(7862):207–212.	1243
1190	Jakob Foerster, Jacob Andreas, Edward Grefenstette,		
1191	Shimon Whiteson, and Tim Rocktäschel. 2019. A	Volodymyr Mnih, Koray Kavukcuoglu, David Silver,	1244
1192	survey of reinforcement learning informed by natural	Andrei A Rusu, Joel Veness, Marc G Bellemare, Alex	1245
1193	language. <i>arXiv preprint arXiv:1906.03926</i> .	Graves, Martin Riedmiller, Andreas K Fidjeland, and	1246
1194	Liangchen Luo, Yinxiao Liu, Rosanne Liu, Sam-	Georg Ostrovski. 2015. Human-level control through	1247
1195	rat Phatale, Meiqi Guo, Harsh Lara, Yunxuan Li,	deep reinforcement learning. <i>Nature</i> , 518(7540):529–	1248
1196	Lei Shu, Yun Zhu, and Lei Meng. 2024. Im-	533.	1249
1197	prove mathematical reasoning in language models	Xuying Ning, Katherine Tieu, Dongqi Fu, Tianxin Wei,	1250
1198	by automated process supervision. <i>arXiv preprint</i>	Zihao Li, Yuanchen Bei, Jiaru Zou, Mengting Ai,	1251
1199	<i>arXiv:2406.06592</i> .	Zhining Liu, and Ting-Wei Li. 2026. Code as agent	1252
1200	Renjie Luo, Zichen Liu, Xiangyan Liu, Chao Du,	harness. <i>arXiv preprint arXiv:2605.18747</i> .	1253
1201	Min Lin, Wenhui Chen, Wei Lu, and Tianyu Pang.		
1202	2025. Language models can learn from verbal	Theo X Olausson, Jeevana Priya Inala, Chenglong	1254
1203	feedback without scalar rewards. <i>arXiv preprint</i>	Wang, Jianfeng Gao, and Armando Solar-Lezama.	1255
1204	<i>arXiv:2509.22638</i> .	2024. Is self-repair a silver bullet for code gener-	1256
1205	Yecheng Jason Ma, William Liang, Guanzhi Wang, De-	ation? In <i>International Conference on Learning</i>	1257
1206	An Huang, Osbert Bastani, Dinesh Jayaraman, Yuke	<i>Representations</i> , volume 2024, pages 36545–36593.	1258
1207	Zhu, and Jim Fan. 2024. Eureka: Human-level re-	Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida,	1259
1208	ward design via coding large language models. In <i>In-</i>	Carroll Wainwright, Pamela Mishkin, Chong Zhang,	1260
1209	<i>ternational conference on learning Representations</i> ,	Sandhini Agarwal, Katarina Slama, and Alex Ray.	1261
1210	volume 2024, pages 26516–26560.	2022. Training language models to follow instruc-	1262
1211	James MacGlashan, Mark K Ho, Robert Loftin, Bei	tions with human feedback. <i>Advances in neural in-</i>	1263
1212	Peng, Guan Wang, David L Roberts, Matthew E	<i>formation processing systems</i> , 35:27730–27744.	1264
1213	Taylor, and Michael L Littman. 2017. Interactive	Liangming Pan, Michael Saxon, Wenda Xu, Deepak	1265
1214	learning from policy-dependent human feedback . In	Nathani, Xinyi Wang, and William Yang Wang. 2024.	1266
1215	<i>International conference on machine learning</i> , pages	Automatically correcting large language models: Sur-	1267
1216	2285–2294. PMLR.	veying the landscape of diverse self-correction strate-	1268
1217	Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler	gies. <i>Transactions of the Association for Computa-</i>	1269
1218	Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon,	<i>tional Linguistics</i> , 12. ArXiv:2308.03188.	1270
1219	Nouha Dziri, Shrimai Prabhumoye, Yiming Yang,	Joon Sung Park, Joseph O’Brien, Carrie Jun Cai, Mered-	1271
1220	Shashank Gupta, Bodhisattwa Prasad Majumder,	ith Ringel Morris, Percy Liang, and Michael S Bern-	1272
1221	Katherine Hermann, Sean Welleck, Amir Yazdan-	stein. 2023. Generative agents: Interactive simulacra	1273
1222	bakhsh, and Peter Clark. 2023. Self-refine: Iterative	of human behavior. In <i>ACM Symposium on User</i>	1274
1223	refinement with self-feedback. In <i>Advances in Neu-</i>	<i>Interface Software and Technology (UIST)</i> .	1275
1224	<i>ral Information Processing Systems</i> .	Debjit Paul, Mete Ismayilzada, Maxime Peyrard, Beat-	1276
1225	Nat McAleese, Rai Michael Pokorny, Juan Felipe Ceron	riz Borges, Antoine Bosselut, Robert West, and	1277
1226	Uribe, Evgenia Nitishinskaya, Maja Trebacz, and	Boi Faltings. 2024. REFINER: Reasoning feed-	1278
1227	Jan Leike. 2024. LLM critics help catch LLM bugs.	back on intermediate representations. <i>arXiv preprint</i>	1279
1228	<i>arXiv preprint arXiv:2407.00215</i> .	<i>arXiv:2304.01904</i> .	1280
1229	Sabrina McCallum, Max Taylor-Davies, Stefano Al-	Moschoula Pternea, Prerna Singh, Abir Chakraborty,	1281
1230	brecht, and Alessandro Suglia. 2023. Is feedback all	Yagna Oruganti, Mirco Milletari, Sayli Bapat, and	1282
1231	you need? leveraging natural language feedback in	Kebei Jiang. 2024. The rl/llm taxonomy tree: Re-	1283
1232	goal-conditioned reinforcement learning . In <i>NeurIPS</i>	viewing synergies between reinforcement learning	1284
1233	<i>2023 Workshop on Goal-Conditioned Reinforcement</i>	and large language models. <i>Journal of Artificial In-</i>	1285
1234	<i>Learning</i> .	<i>telligence Research</i> , 80:1525–1573.	1286
		Keyang Qian, Yixin Cheng, Rui Guan, Wei Dai, Flora	1287
		Jin, Kaixun Yang, Sadia Nawaz, Zachari Swiecki,	1288
		Guanliang Chen, and Lixiang Yan. 2025. Dean of	1289
		LLM tutors: A framework for automated quality	1290

1291	review of AI-generated feedback. <i>arXiv preprint arXiv:2508.05952</i> .	1347
1292		1348
1293	Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D Manning, Stefano Ermon, and Chelsea Finn. 2023. Direct preference optimization: Your language model is secretly a reward model. <i>Advances in neural information processing systems</i> , 36:53728–53741.	1349
1294		1350
1295		1351
1296		
1297		
1298	Yusuke Sakai, Hidetaka Kamigaito, and Taro Watanabe. 2025. Revisiting compositional generalization capability of large language models considering instruction following ability. In <i>Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 31219–31238.	1352
1299		1353
1300		1354
1301		1355
1302		1356
1303		1357
1304		
1305	Jérémy Scheurer, Jon Ander Campos, Tomasz Korbak, Jun Shern Chan, Angelica Chen, Kyunghyun Cho, and Ethan Perez. 2023. Training language models with language feedback at scale. <i>arXiv preprint arXiv:2303.16755</i> .	1358
1306		1359
1307		1360
1308		1361
1309		1362
1310		1363
1311		
1312	Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Eric Hambro, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. 2023. Toolformer: Language models can teach themselves to use tools. <i>Advances in neural information processing systems</i> , 36:68539–68551.	1364
1313		1365
1314		1366
1315		1367
1316		
1317	Mrinank Sharma, Meg Tong, Tomek Korbak, David Duvenaud, Amanda Askell, Sam Bowman, Esin Durmus, Zac Hatfield-Dodds, Scott Johnston, and Shauna Kravec. 2024. Towards understanding sycophancy in language models. In <i>International Conference on Learning Representations</i> , volume 2024, pages 110–144.	1368
1318		1369
1319		1370
1320		1371
1321		1372
1322		
1323	Pratyusha Sharma, Balakumar Sundaralingam, Valts Blukis, Chris Paxton, Tucker Hermans, Antonio Torralba, Jacob Andreas, and Dieter Fox. 2022. Correcting robot plans with natural language feedback. <i>arXiv preprint arXiv:2204.05186</i> .	1373
1324		1374
1325		1375
1326		1376
1327		
1328	Haizhou Shi, Ye Liu, Bo Pang, Zeyu Leo Liu, Hao Wang, Silvio Savarese, Caiming Xiong, Yingbo Zhou, and Semih Yavuz. 2025a. Ssr: Socratic self-refine for large language model reasoning. <i>arXiv preprint arXiv:2511.10621</i> .	1377
1329		1378
1330		1379
1331		1380
1332		1381
1333		
1334	Jiawen Shi, Zenghui Yuan, Guiyao Tie, Pan Zhou, Neil Zhenqiang Gong, and Lichao Sun. 2025b. Prompt injection attack to tool selection in llm agents. <i>arXiv preprint arXiv:2504.19793</i> .	1382
1335		1383
1336		1384
1337		1385
1338	Noah Shinn, Federico Cassano, Berman Labash, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. 2023. Reflexion: Language agents with verbal reinforcement learning. In <i>Advances in Neural Information Processing Systems</i> , volume 36.	1386
1339		1387
1340		1388
1341		1389
1342		
1343	Mohit Shridhar, Xingdi Yuan, Marc-Alexandre Côté, Yonatan Bisk, Adam Trischler, and Matthew Hausknecht. 2020. Alfworld: Aligning text and embodied environments for interactive learning. <i>arXiv preprint arXiv:2010.03768</i> .	1390
1344		1391
1345		1392
1346		1393
		1394
		1395
	Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Yarin Gal, Nicolas Papernot, and Ross Anderson. 2023. The curse of recursion: Training on generated data makes models forget. <i>arXiv preprint arXiv:2305.17493</i> .	1396
		1397
		1398
		1399
		1400
	David Silver, Aja Huang, Chris J Maddison, Arthur Guez, Laurent Sifre, George Van Den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, and Marc Lanctot. 2016. Mastering the game of go with deep neural networks and tree search. <i>nature</i> , 529(7587):484–489.	
	Avi Singh, John D Co-Reyes, Rishabh Agarwal, Ankesh Anand, Piyush Patil, Xavier Garcia, Peter J Liu, James Harrison, Jaehoon Lee, and Kelvin Xu. 2023. Beyond human data: Scaling self-training for problem-solving with language models. <i>arXiv preprint arXiv:2312.06585</i> .	
	Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. 2024. Scaling LLM test-time compute optimally can be more effective than scaling model parameters. <i>arXiv preprint arXiv:2408.03314</i> .	
	Yuda Song, Lili Chen, Fahim Tajwar, Remi Munos, Deepak Pathak, J Andrew Bagnell, Aarti Singh, and Andrea Zanette. 2026. Expanding the capabilities of reinforcement learning via text feedback. <i>arXiv preprint arXiv:2602.02482</i> .	
	Alexander L Strehl, Lihong Li, and Michael L Littman. 2009. Reinforcement learning in finite MDPs: PAC analysis. <i>Journal of Machine Learning Research</i> , 10(84):2413–2444.	
	Shengjie Sun, Runze Liu, Jiafei Lyu, Jing-Wen Yang, Liangpeng Zhang, and Xiu Li. 2025a. A large language model-driven reward design framework via dynamic feedback for reinforcement learning. <i>Knowledge-Based Systems</i> , 326:114065.	
	Shengjie Sun, Jiafei Lyu, Runze Liu, Mengbei Yan, Bo Liu, Deheng Ye, and Xiu Li. 2025b. Prof: An llm-based reward code preference optimization framework for offline imitation learning. <i>arXiv preprint arXiv:2511.13765</i> .	
	Richard S Sutton and Andrew G Barto. 1998. <i>Reinforcement learning: An introduction</i> , volume 1. MIT press Cambridge.	
	Haoran Tan, Zeyu Zhang, Chen Ma, Xu Chen, Quanyu Dai, and Zhenhua Dong. 2025. Membench: Towards more comprehensive evaluation on the memory of llm-based agents. In <i>Findings of the Association for Computational Linguistics: ACL 2025</i> , pages 19336–19352.	
	Jonathan Uesato, Nate Kushman, Ramana Kumar, Francis Song, Noah Siegel, Lisa Wang, Antonia Creswell, Geoffrey Irving, and Irina Higgins. 2022. Solving math word problems with process-and outcome-based feedback. <i>arXiv preprint arXiv:2211.14275</i> .	

1401	Peter Vamplew, Benjamin J. Smith, Johan Kallstrom,	Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten	1458
1402	Gabriel Ramos, Roxana Radulescu, Diederik M.	Bosma, Fei Xia, Ed Chi, Quoc V Le, and Denny	1459
1403	Roijsers, Conor F. Hayes, Fredrik Heintz, Patrick	Zhou. 2022. Chain-of-thought prompting elicits rea-	1460
1404	Mannion, Pieter J. K. Libin, Richard Dazeley, and	soning in large language models. <i>Advances in neural</i>	1461
1405	Cameron Foale. 2021. Scalar reward is not enough:	<i>information processing systems</i> , 35:24824–24837.	1462
1406	A response to silver, singh, precup and sutton (2021).		
1407	<i>Preprint</i> , arXiv:2112.15422.		
1408	Eric Wallace, Kai Xiao, Reimar Leike, Lilian Weng,	Wenyi Xiao, Zechuan Wang, Leilei Gan, Shuai Zhao,	1463
1409	Johannes Heidecke, and Alex Beutel. 2024. The in-	Zongrui Li, Ruirui Lei, and 1 others. 2024. A com-	1464
1410	struction hierarchy: Training llms to prioritize privi-	prehensive survey of Direct Preference Optimization:	1465
1411	leged instructions. <i>arXiv preprint arXiv:2404.13208</i> .	Datasets, theories, variants, and applications. <i>arXiv</i>	1466
		<i>preprint arXiv:2410.15595</i> .	1467
1412	Ante Wang, Linfeng Song, Ye Tian, Dian Yu, Haitao	Tianbao Xie, Siheng Zhao, Chen Wu, Yitao Liu, Qian	1468
1413	Mi, Xiangyu Duan, Zhaopeng Tu, Jinsong Su, and	Luo, Victor Zhong, Yanchao Yang, and Tao Yu. 2024.	1469
1414	Dong Yu. 2025a. Don’t get lost in the trees: Stream-	Text2reward: Reward shaping with language models	1470
1415	lining LLM reasoning by overcoming tree search	for reinforcement learning. In <i>International Con-</i>	1471
1416	exploration pitfalls . In <i>Proceedings of the 63rd An-</i>	<i>ference on Learning Representations</i> , volume 2024,	1472
1417	<i>annual Meeting of the Association for Computational</i>	pages 35663–35699.	1473
1418	<i>Linguistics (Volume 1: Long Papers)</i> , pages 23946–		
1419	23959, Vienna, Austria. Association for Computa-	Yuxi Xie, Kenji Kawaguchi, Yiran Zhao, James Xu	1474
1420	tional Linguistics.	Zhao, Min-Yen Kan, Junxian He, and Michael Xie.	1475
		2023. Self-evaluation guided beam search for rea-	1476
1421	Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Man-	soning. <i>Advances in Neural Information Processing</i>	1477
1422	dlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and An-	<i>Systems</i> , 36:41618–41650.	1478
1423	ima Anandkumar. 2023a. Voyager: An open-ended		
1424	embodied agent with large language models. <i>arXiv</i>	John Yang, Carlos E Jimenez, Alexander Wettig, Kilian	1479
1425	<i>preprint arXiv:2305.16291</i> .	Lieret, Shunyu Yao, Karthik Narasimhan, and Ofir	1480
		Press. 2024. Swe-agent: Agent-computer interfaces	1481
1426	Hanyang Wang, Lu Wang, Chaoyun Zhang, Tianjun	enable automated software engineering. <i>Advances in</i>	1482
1427	Mao, Si Qin, Qingwei Lin, Saravan Rajmohan, and	<i>Neural Information Processing Systems</i> , 37:50528–	1483
1428	Dongmei Zhang. 2025b. Text2grad: Reinforcement	50652.	1484
1429	learning from natural language feedback. <i>arXiv</i>		
1430	<i>preprint arXiv:2505.22338</i> .	Shunyu Yao, Rohan Rao, Matthew Hausknecht, and	1485
		Karthik Narasimhan. 2020. Keep calm and explore:	1486
1431	Peiyi Wang, Lei Li, Zhihong Shao, R. X. Xu, Damai	Language models for action generation in text-based	1487
1432	Dai, Yifei Li, Deli Chen, Y. Wu, and Zhifang Sui.	games. In <i>Proceedings of the 2020 Conference on</i>	1488
1433	2024a. Math-shepherd: Verify and reinforce llms	<i>Empirical Methods in Natural Language Processing</i>	1489
1434	step-by-step without human annotations . <i>Preprint</i> ,	(EMNLP), pages 8736–8754.	1490
1435	arXiv:2312.08935.		
1436	Qibin Wang, Pu Zhao, Shaohan Huang, Fangkai Yang,	Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran,	1491
1437	Lu Wang, Furu Wei, Qingwei Lin, Saravan Rajmo-	Tom Griffiths, Yuan Cao, and Karthik Narasimhan.	1492
1438	han, and Dongmei Zhang. 2025c. Learning to refine:	2023. Tree of thoughts: Deliberate problem solving	1493
1439	Self-refinement of parallel reasoning in llms. <i>arXiv</i>	with large language models. <i>Advances in neural</i>	1494
1440	<i>preprint arXiv:2509.00084</i> .	<i>information processing systems</i> , 36:11809–11822.	1495
1441	Tianlu Wang, Ping Yu, Xiaoqing Ellen Tan, Sean	Weizhe Yuan, Richard Yuanzhe Pang, Kyunghyun Cho,	1496
1442	O’Brien, Ramakanth Pasunuru, Jane Dwivedi-Yu,	Xian Li, Sainbayar Sukhbaatar, Jing Xu, and Ja-	1497
1443	Olga Golovneva, Luke Zettlemoyer, Maryam Fazel-	son Weston. 2024. Self-rewarding language models.	1498
1444	Zarandi, and Asli Celikyilmaz. 2023b. Shepherd: A	<i>arXiv preprint arXiv:2401.10020</i> .	1499
1445	critic for language model generation. <i>arXiv preprint</i>		
1446	<i>arXiv:2308.04592</i> .	Eric Zelikman, Yuhuai Wu, Jesse Mu, and Noah D	1500
1447	Xingyao Wang, Yangyi Chen, Lifan Yuan, Yizhe Zhang,	Goodman. 2022. STaR: Bootstrapping reasoning	1501
1448	Yunzhu Li, Hao Peng, and Heng Ji. 2024b. Exe-	with reasoning. In <i>Advances in Neural Information</i>	1502
1449	cutable code actions elicit better llm agents. In <i>Forty-</i>	<i>Processing Systems</i> .	1503
1450	<i>first International Conference on Machine Learning</i> .		
1451	Xingyao Wang, Boxuan Li, Yufan Song, Frank F Xu,	Runhao Zeng, Dingjie Zhou, Qiwei Liang, Junlin Liu,	1504
1452	Xiangru Tang, Mingchen Zhuge, Jiayi Pan, Yueqi	Hui Li, Changxin Huang, Jianqiang Li, Xiping Hu,	1505
1453	Song, Bowen Li, and Jaskirat Singh. 2025d. Open-	and Fuchun Sun. 2024. Video2reward: Generating	1506
1454	hands: An open platform for ai software developers	reward function from videos for legged robot behav-	1507
1455	as generalist agents. In <i>International Conference</i>	ior learning. <i>arXiv preprint arXiv:2412.05515</i> .	1508
1456	<i>on Learning Representations</i> , volume 2025, pages		
1457	65882–65919.	Qiusi Zhan, Richard Fang, Henil Shalin Panchal, and	1509
		Daniel Kang. 2025. Adaptive attacks break de-	1510
		fenses against indirect prompt injection attacks on	1511
		llm agents. In <i>Findings of the Association for Com-</i>	1512
		<i>putational Linguistics: NAACL 2025</i> , pages 7101–	1513
		7117.	1514

- Qiusi Zhan, Zhixiang Liang, Zifan Ying, and Daniel Kang. 2024. Injecagent: Benchmarking indirect prompt injections in tool-integrated large language model agents. In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 10471–10506.
- Guibin Zhang, Hejia Geng, Xiaohang Yu, Zhenfei Yin, Zaibin Zhang, Zelin Tan, and 1 others. 2025a. The landscape of agentic reinforcement learning for LLMs: A survey. *arXiv preprint arXiv:2509.02547*.
- Qiyuan Zhang, Fuyuan Lyu, Zexu Sun, Lei Wang, Weixu Zhang, Wenyue Hua, and 1 others. 2025b. A survey on test-time scaling in large language models: What, how, where, and how well? *arXiv preprint arXiv:2503.24235*.
- Xiaohan Zhang, Yan Ding, Yohei Hayamizu, Zainab Altaweel, Yifeng Zhu, Yuke Zhu, Peter Stone, Chris Paxton, and Shiqi Zhang. 2025c. Llm-grop: Visually grounded robot task and motion planning with large language models. *The International Journal of Robotics Research*, page 02783649251378196.
- Yang Zhang, Shixin Yang, Chenjia Bai, Fei Wu, Xiu Li, Zhen Wang, and Xuelong Li. 2025d. Towards efficient llm grounding for embodied multi-agent collaboration. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 1663–1699.
- Zeyu Zhang, Quanyu Dai, Xiaohe Bo, Chen Ma, Rui Li, Xu Chen, Jieming Zhu, Zhenhua Dong, and Ji-Rong Wen. 2025e. A survey on the memory mechanism of large language model-based agents. *ACM Transactions on Information Systems*, 43(6):1–47.
- Andrew Zhao, Daniel Huang, Quentin Xu, Matthieu Lin, Yong-Jin Liu, and Gao Huang. 2024. ExpeL: LLM agents are experiential learners. In *AAAI Conference on Artificial Intelligence*.
- Jian Zhao, Runze Liu, Kaiyan Zhang, Zhimu Zhou, Junqi Gao, Dong Li, Jiafei Lyu, Zhouyi Qian, Biqing Qi, and Xiu Li. 2026. Genprm: Scaling test-time compute of process reward models via generative reasoning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pages 34932–34940.
- Chujie Zheng, Zhenru Zhang, Beichen Zhang, Runji Lin, Keming Lu, Bowen Yu, and 1 others. 2024. ProcessBench: Identifying process errors in mathematical reasoning. *arXiv preprint arXiv:2412.06559*.
- Enshen Zhou, Qi Su, Cheng Chi, Zhizheng Zhang, Zhongyuan Wang, Tiejun Huang, Lu Sheng, and He Wang. 2025a. Code-as-monitor: Constraint-aware visual programming for reactive and proactive robotic failure detection. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pages 6919–6929.
- Yan Zhou and Yanguang Chen. 2025. Adaptive heterogeneous multi-agent debate for enhanced educational and factual reasoning in large language models. *Journal of King Saud University Computer and Information Sciences*, 37(10):330.
- Yuxuan Zhou, Yubin Wang, Bin Wang, Chen Ning, Xien Liu, Ji Wu, and Jianye Hao. 2025b. Enhancing the medical context-awareness ability of llms via multifaceted self-refinement learning. *arXiv preprint arXiv:2511.10067*.

A Comparison with Prior Surveys

This survey is, to our knowledge, the first organized around the verbal signal itself: which component of the decision problem a linguistic artifact modifies and when it takes effect. Prior overviews organize by mechanism (how a method works) or by agent capability (what a system does). Table 3 makes the positioning explicit. Relative to each row, our contribution is the signal-centric axis, the inclusion criteria of Section 1, the hybrid-case decision rules of Table 1, and the practitioner guidance of Appendix B.

B Choosing a Verbal-Feedback Mechanism

Table 4 distills the survey into prescriptive guidance. Each row states when a mechanism is the right choice, keyed to properties a practitioner can assess before building: whether task outcomes are verifiable, whether model weights are accessible, the latency and cost budget for feedback, and the dominant risk to monitor. Representative methods are drawn from the corresponding sections.

Table 3: Positioning against prior surveys and frameworks. Axis: the work’s primary organizing principle. Feedback: signal types covered (S = scalar, P = preference, V = verbal). Stages: agent-lifecycle stages covered (D = problem definition, I = inference, T = training). Guide: whether prescriptive practitioner guidance is offered.

Survey / framework	Axis	Feedback	Stages	Guide
Luketina et al. (2019)	Language-conditional vs. language-assisted RL	S	D, T	✗
Pan et al. (2024)	When correction happens	V	I, T	✗
Kamoi et al. (2024)	When self-correction works	V	I	✗
Fernandes et al. (2023)	Feedback format and usage for generation	S, P, V	T	✗
Kaufmann et al. (2023)	RLHF pipeline	S, P	T	✗
Xiao et al. (2024)	Preference-optimization variants	P	T	✗
Zhang et al. (2025e)	Agent memory mechanisms	V	I	✗
Zhang et al. (2025b)	Test-time scaling (what/how/where)	S, V	I	✗
Gao et al. (2025)	What/when/how agents evolve	S, V	I, T	✗
Fang et al. (2025)	Agent-system optimizers	S, V	I, T	✗
Zhang et al. (2025a)	RL for agent capabilities	S	T	✗
Feng et al. (2024)	RL components in language space (framework)	V	I, T	✗
This survey	The verbal signal: what it modifies, when it acts	V (S, P as boundary)	D, I, T	✓

Table 4: Practitioner guidance: when to use language in each role.

Use language as	When	Dominant risk	Representative
Reward (compile to code or rubric)	Success is expressible as checks; simulator or executor available; training budget exists	Objective mis-specification; reward hacking	Eureka (Ma et al., 2024); Text2Reward (Xie et al., 2024)
State / task grounding	Raw observations too complex or task underspecified; language abstracts them faithfully	Information loss; state aliasing	ALFWorld (Shridhar et al., 2020); SayCan (Ahn et al., 2022)
Deliberative feedback	No weight access; latency budget allows iteration; a grounded critic (tool, tests) exists	Circular self-critique; inference cost	Self-Refine (Madaan et al., 2023); CRITIC (Gou et al., 2024a); ToT (Yao et al., 2023)
Memory	Tasks recur across episodes; lessons are reusable	Error persistence; memory poisoning	Reflexion (Shinn et al., 2023); ExpeL (Zhao et al., 2024); Voyager (Wang et al., 2023a)
Training signal	Persistent improvement needed; feedback plentiful; weights accessible	Sycophancy to wrong critiques; filter bias	FCP (Luo et al., 2025); STaR (Zelikman et al., 2022); PRMs (Lightman et al., 2024)